

Spectral clustering

Reading https://people.csail.mit.edu/dsontag/courses/ml14/notes/Luxburg07_tutorial_spectral_clustering.pdf

Spectral clustering

Steps

- Preprocess the data
- Find the graph Laplacian vector of the matrix
- Use K-means on eigenvectors of the graph Laplacian to cluster

First, compute AX for

$$X = \begin{bmatrix} -5 \\ -4 \\ 3 \end{bmatrix}$$

This product is given by

$$AX = \begin{bmatrix} 0 & 5 & -10 \\ 0 & 22 & 16 \\ 0 & -9 & -2 \end{bmatrix} \begin{bmatrix} -5 \\ -4 \\ 3 \end{bmatrix} = \begin{bmatrix} -50 \\ -40 \\ 30 \end{bmatrix} = 10 \begin{bmatrix} -5 \\ -4 \\ 3 \end{bmatrix}$$

In this case, the product AX resulted in a vector which is equal to 10 times the vector X . In other words, $AX = 10X$.

Perhaps this is a special matrix A which when multiplied by any vector \mathbf{x} always equals $k\mathbf{x}$?

$$\begin{bmatrix} 0 & 5 & -10 \\ 0 & 22 & 16 \\ 0 & -9 & -2 \end{bmatrix} \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix} = \begin{bmatrix} -5 \\ 38 \\ -11 \end{bmatrix}$$

In this case, AX did not result in a vector of the form kX for some scalar k .

NOPE!

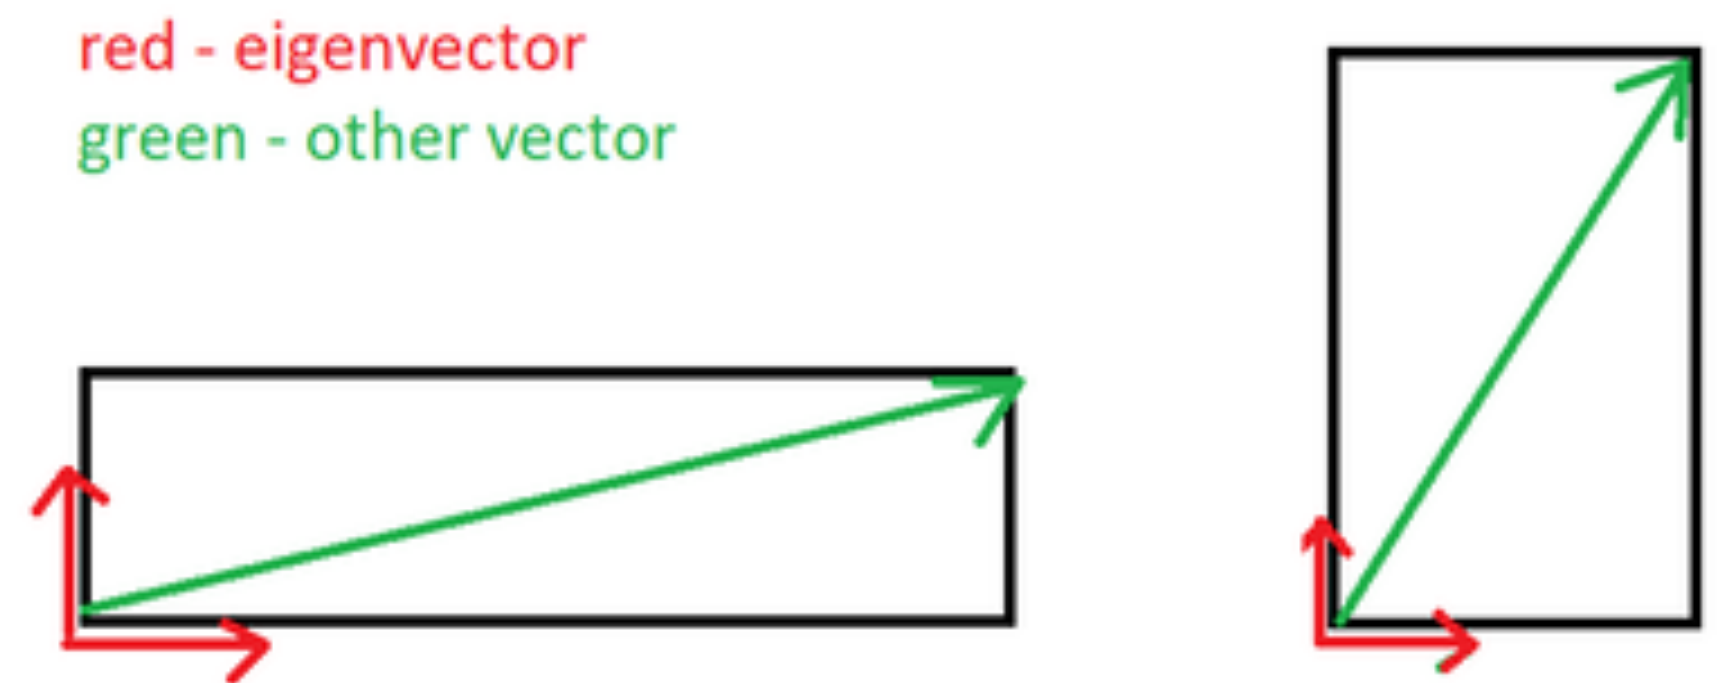
Eigenvectors and Eigenvalues

When $A \in \mathbb{R}^{n \times n}$ is positive semidefinite then a vector \mathbf{v} is an eigenvector and has associated eigenvalue λ iff

$$(A - \lambda I)\mathbf{v} = 0$$

$$A\mathbf{v} = \lambda\mathbf{v}$$

That is eigenvectors do not change direction, they only scale linearly when A is subject to a linear transformation.




Eigenvectors and Eigenvalues

If $A \in \mathbb{R}^{n \times n}$ is symmetric then there exists an orthogonal set of eigenvectors.

Define $U = [\mathbf{v}_1, \mathbf{v}_2 \dots \mathbf{v}_n]$, a square matrix whose columns are the n linearly independent eigenvectors of \mathbf{A} . Also define the vector of matching eigenvalues $\sigma(A) = [\lambda_1, \lambda_2 \dots \lambda_n]$; this vector is known as the spectrum of A . Lastly define Λ , the diagonalization of the spectrum $\Lambda = \mathbf{I}\sigma(A)$

It can be shown that $U^{-1}AU = \Lambda$



Eigenvectors and Eigenvalues

Can be generalized to $A \in \mathbb{R}^{n \times m}$ if A is invertible (full rank/positive definite) through Singular Value Decomposition

Let A be an $m \times n$ matrix. Then there exist orthogonal matrices U and V of the appropriate size such that $A = U\Sigma V^T$ where Σ is of the form

$$\Sigma = \begin{bmatrix} \sigma & 0 \\ 0 & 0 \end{bmatrix}$$

and σ is of the form

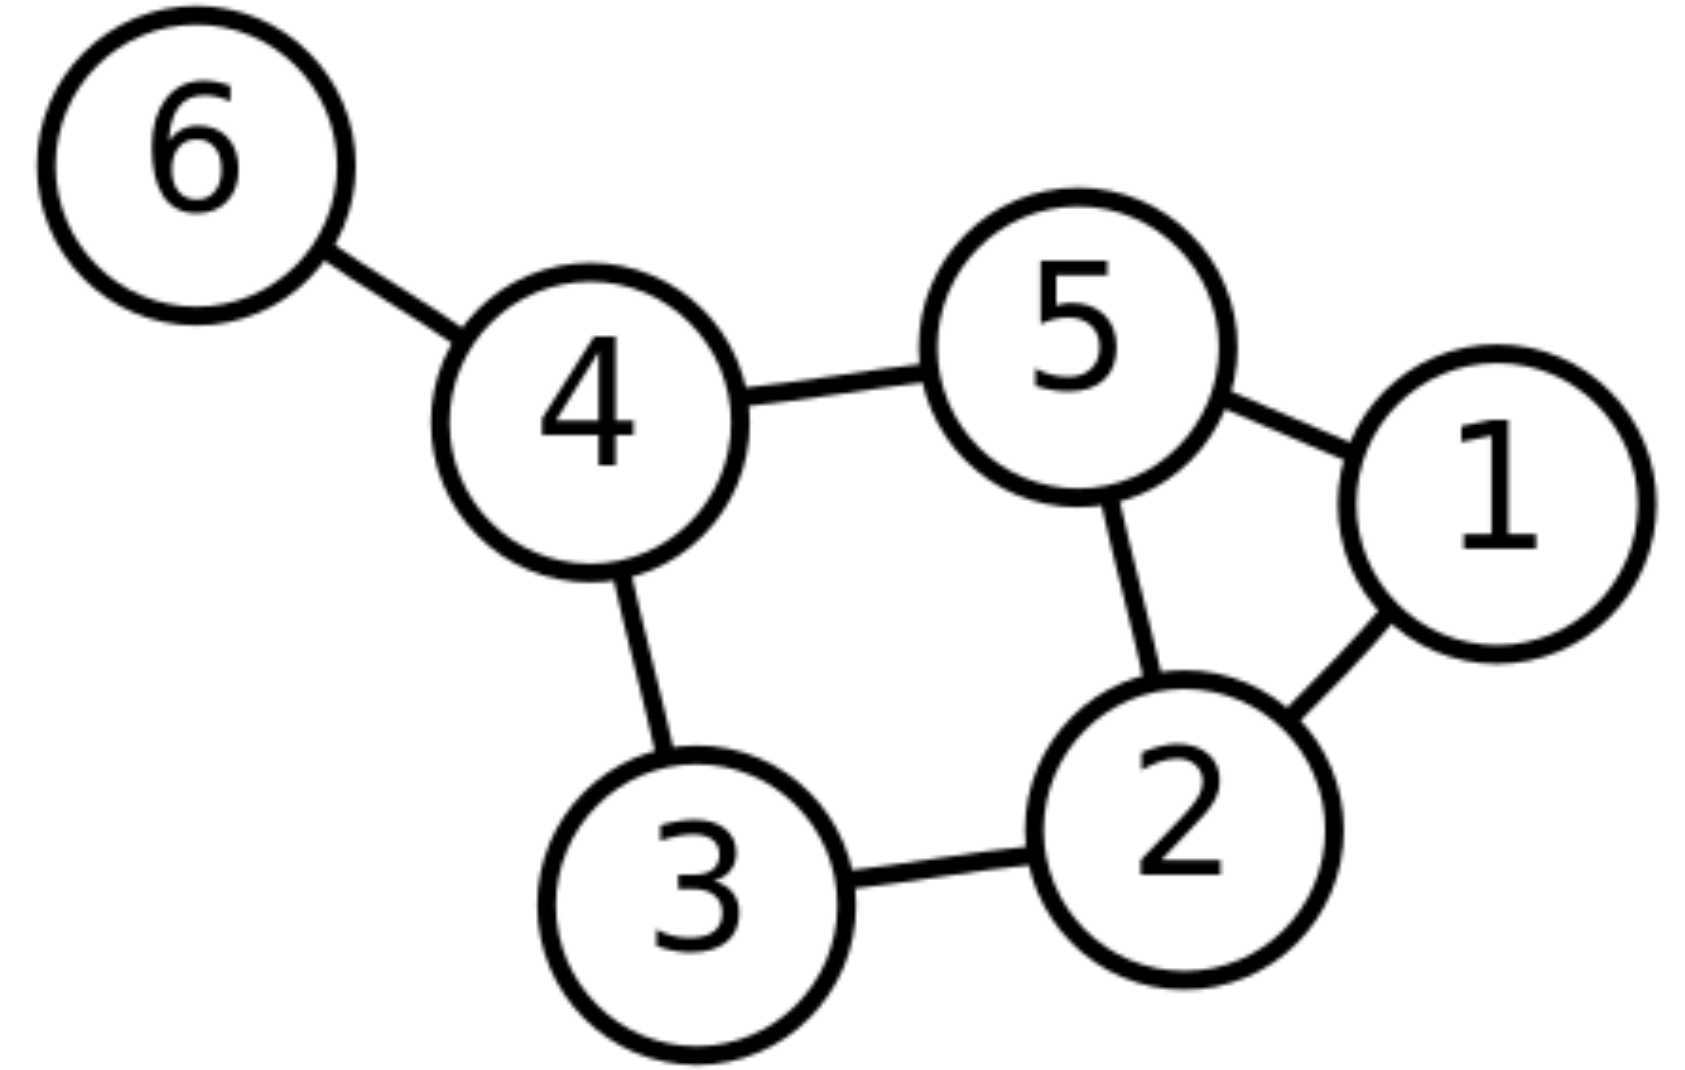
$$\sigma = \begin{bmatrix} \sigma_1 & & 0 \\ & \ddots & \\ 0 & & \sigma_k \end{bmatrix}$$

for the σ_i the singular values of A .

In one restricted but very common sense of the term,^{[1][2]} a **graph** is an **ordered pair** $G = (V, E)$ comprising:

- V , a **set** of **vertices** (also called **nodes** or **points**);
- $E \subseteq \{\{x, y\} \mid x, y \in V \text{ and } x \neq y\}$, a **set** of **edges** (also called **links** or **lines**), which are **unordered pairs** of vertices (that is, an edge is associated with two distinct vertices).


To avoid ambiguity, this type of object may be called precisely an **undirected simple graph**.



Similarity graph

There are many ways to make a graph from data


- Unweighted
 - $W_{i,j} = 1$ iff $d(x_i, x_j) < \epsilon \ \forall i, j \neq i \in G$; otherwise 0
 - $W_{i,j} = 1$ iff x_i is in the set of k nearest neighbors of x_j ; otherwise 0
- Weighted
 - $W_{i,j} = \exp(-\frac{\|x_i - x_j\|^2}{2\sigma^2})$



Affinity matrix

Unweighted example

	1	2	3	4	5	6
1	0	1	0	0	1	0
2	1	0	1	0	1	0
3	0	1	0	1	0	0
4	0	0	1	0	1	1
5	1	1	0	1	0	0
6	0	0	0	1	0	0

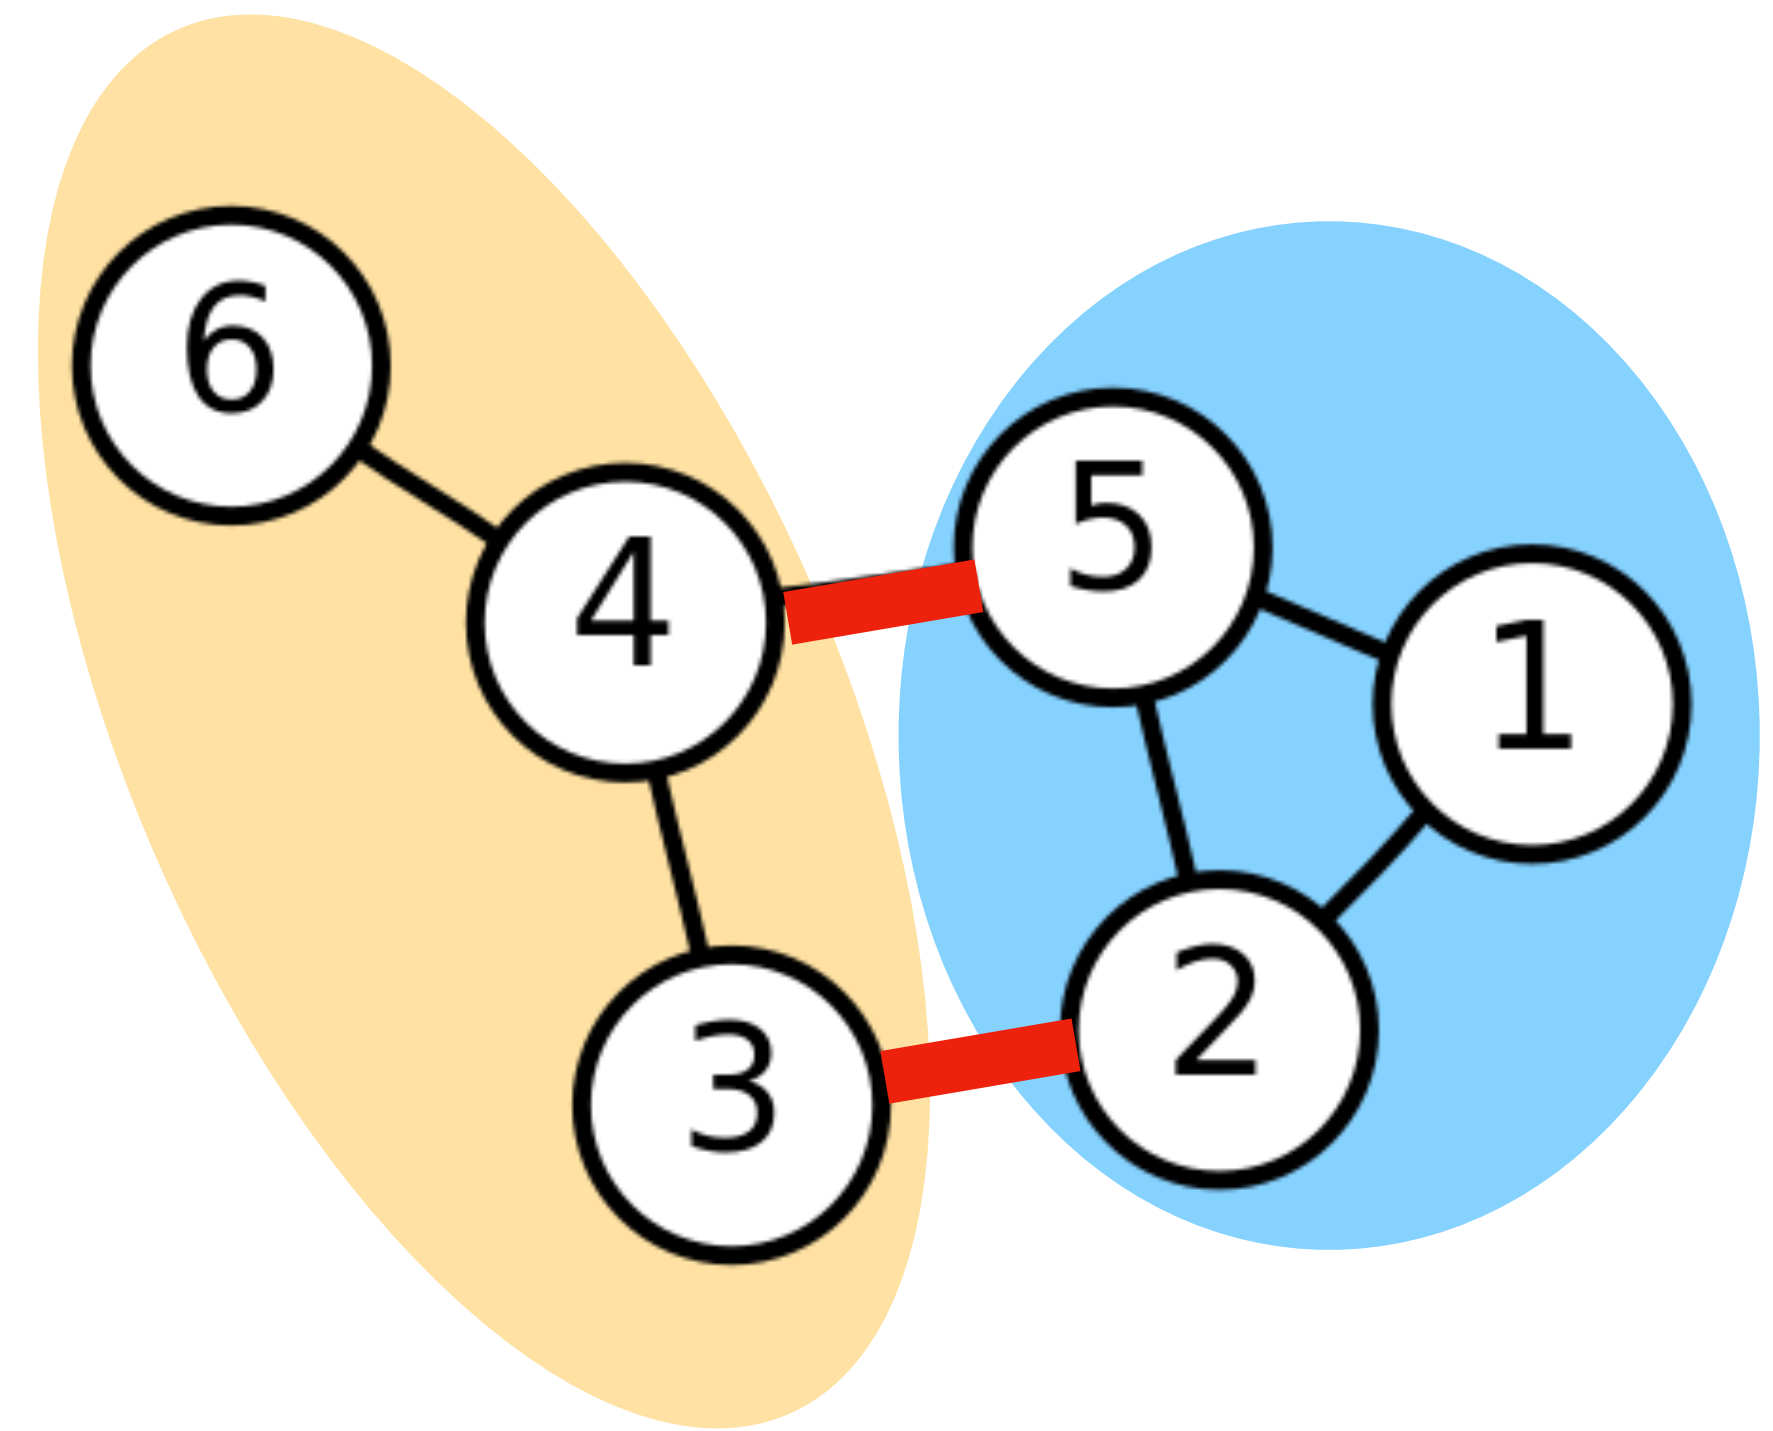


Clustering as a bi-partitioning task

Partition vertices into two disjoint groups A, B

How can we define a “good” partition?

How can we efficiently calculate such a partition?



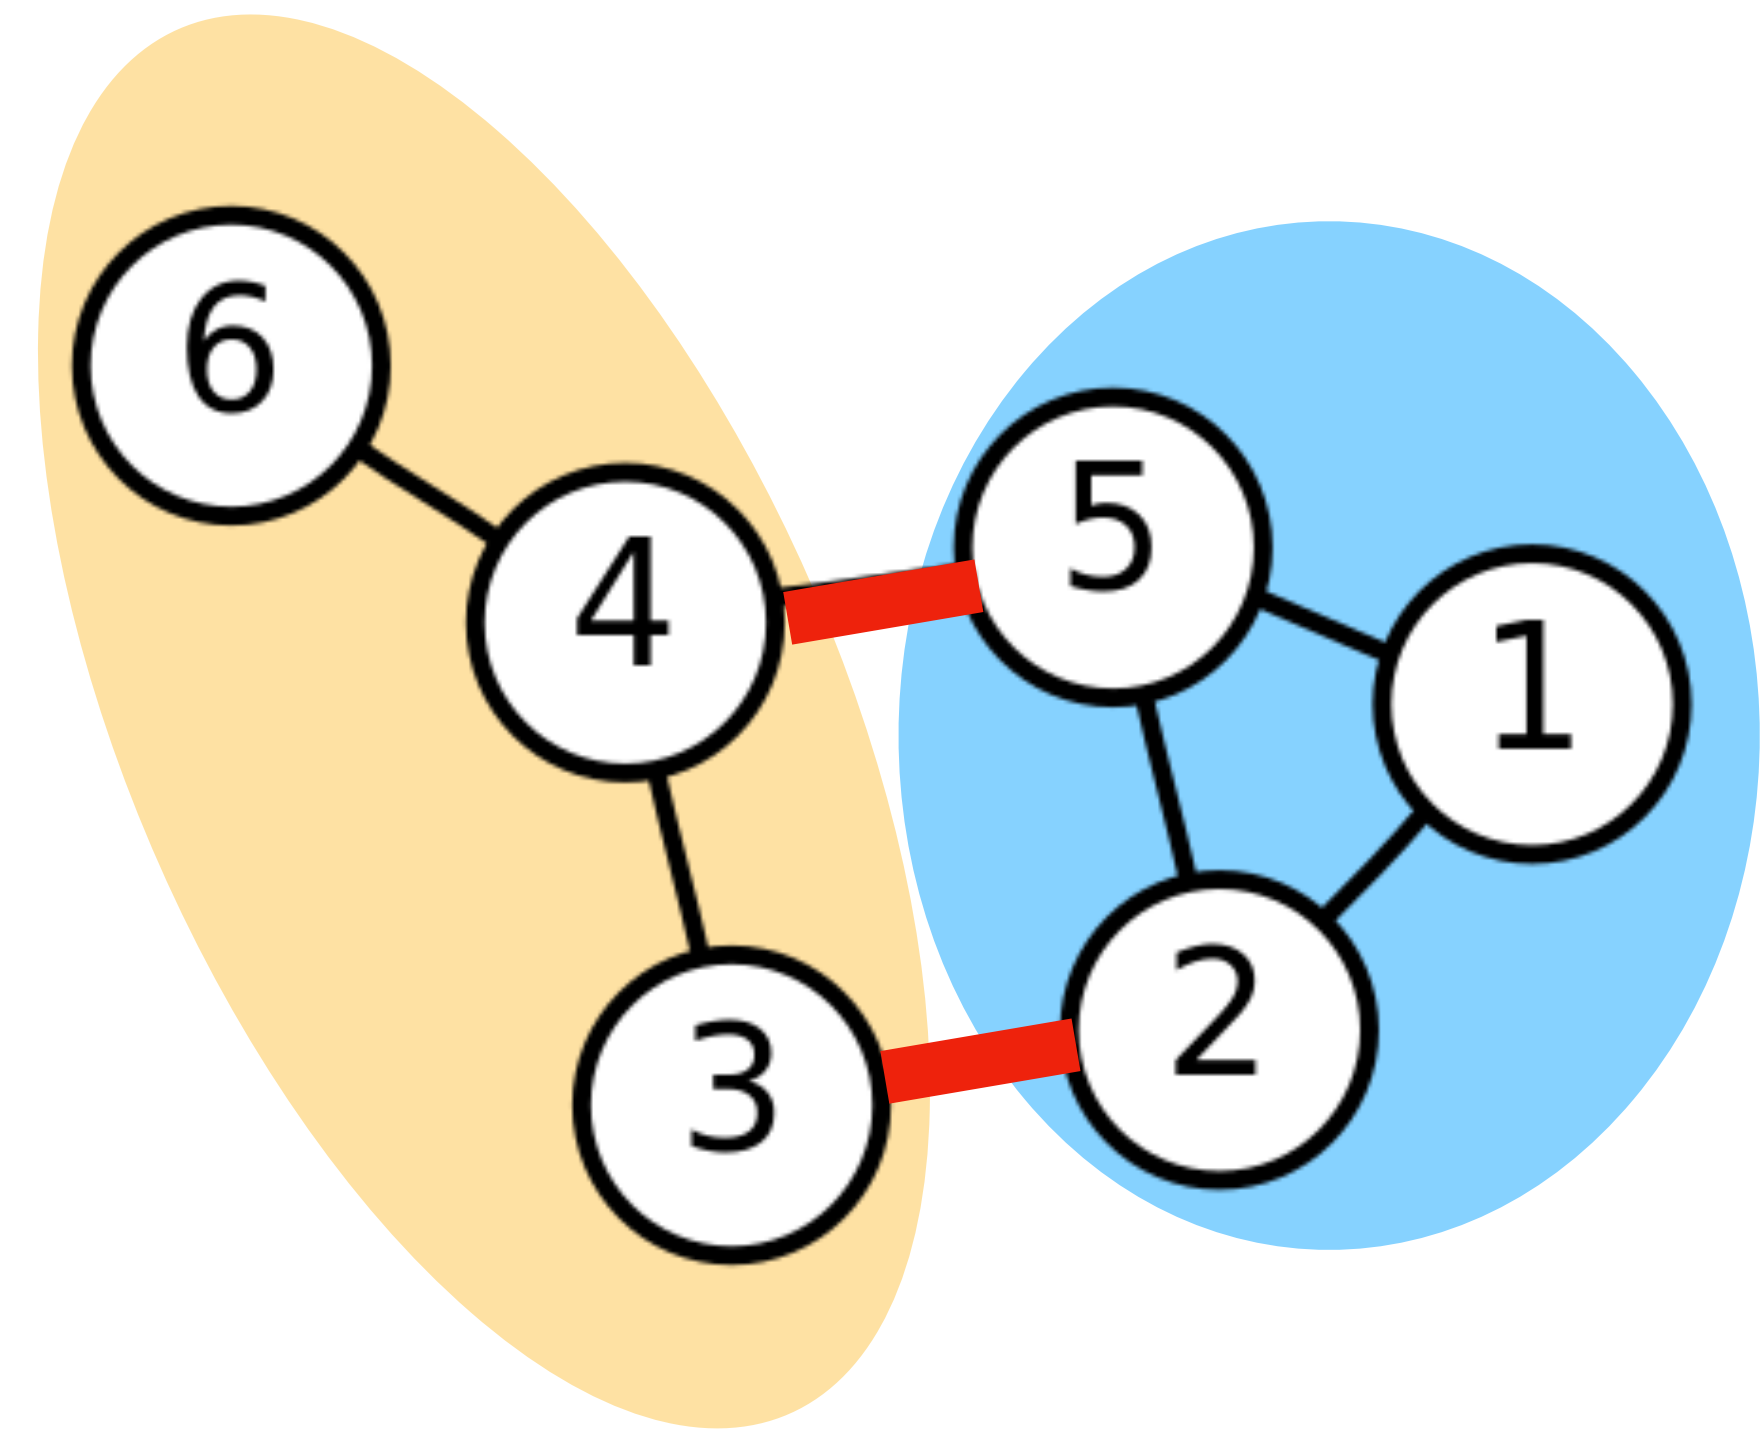
Clustering as a bi-partitioning task

Partition vertices into two disjoint groups A, B

$$\text{cut}(A,B) = \sum_{\forall i \in A, \forall j \in B} W_{i,j}$$

And just

$$\arg \min_{A,B} \text{cut}(A,B)?$$

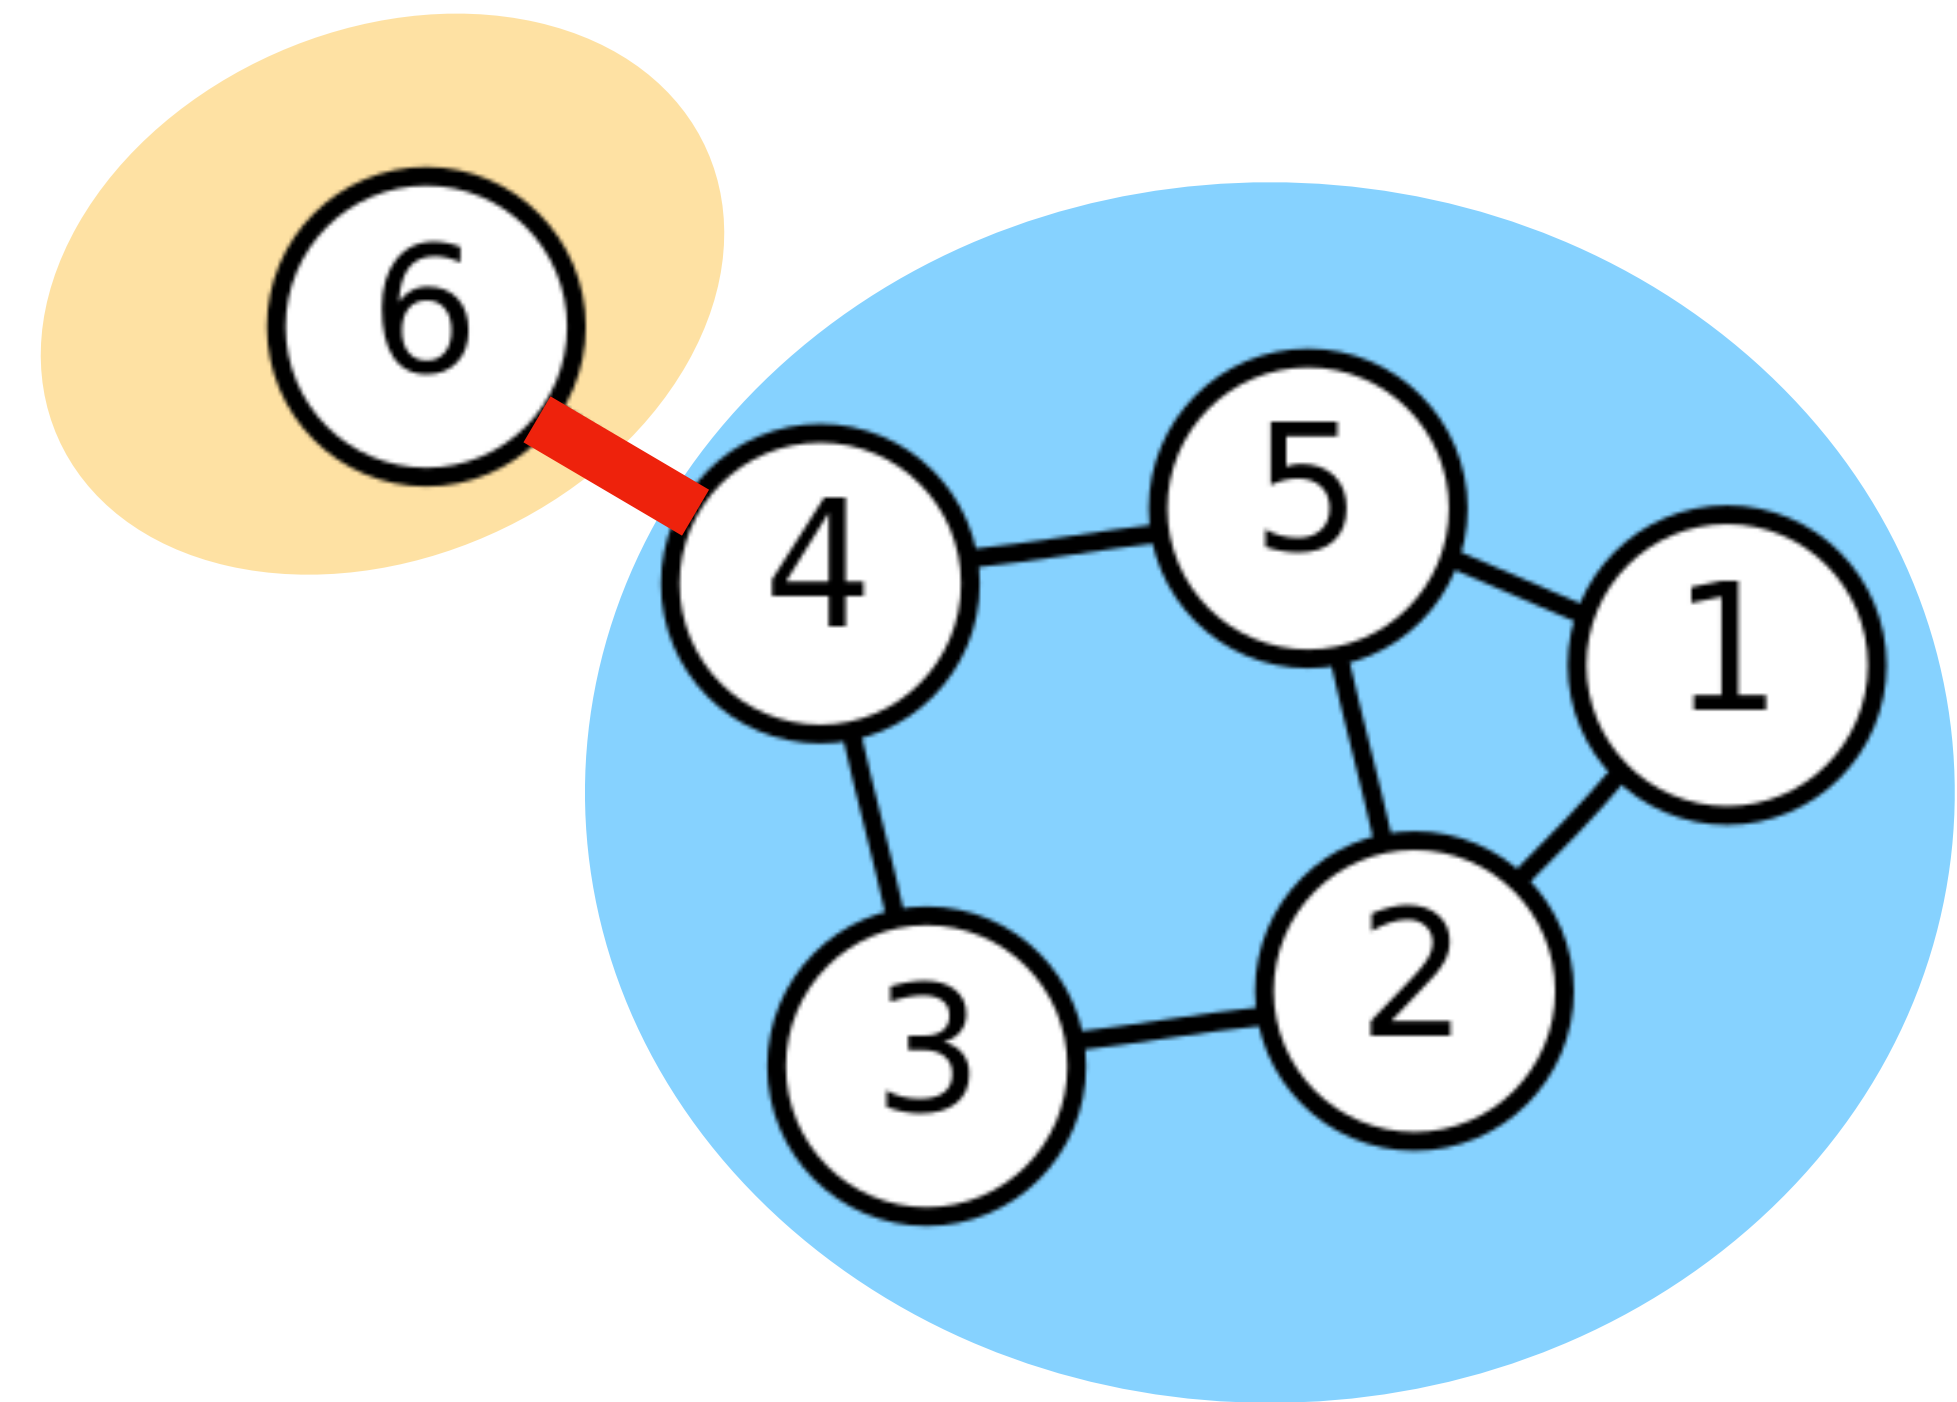


Clustering as a bi-partitioning task

Partition vertices into two disjoint groups A, B

$$\arg \min_{A,B} \text{cut}(A,B)$$

Degenerate case!



Clustering as a bi-partitioning task

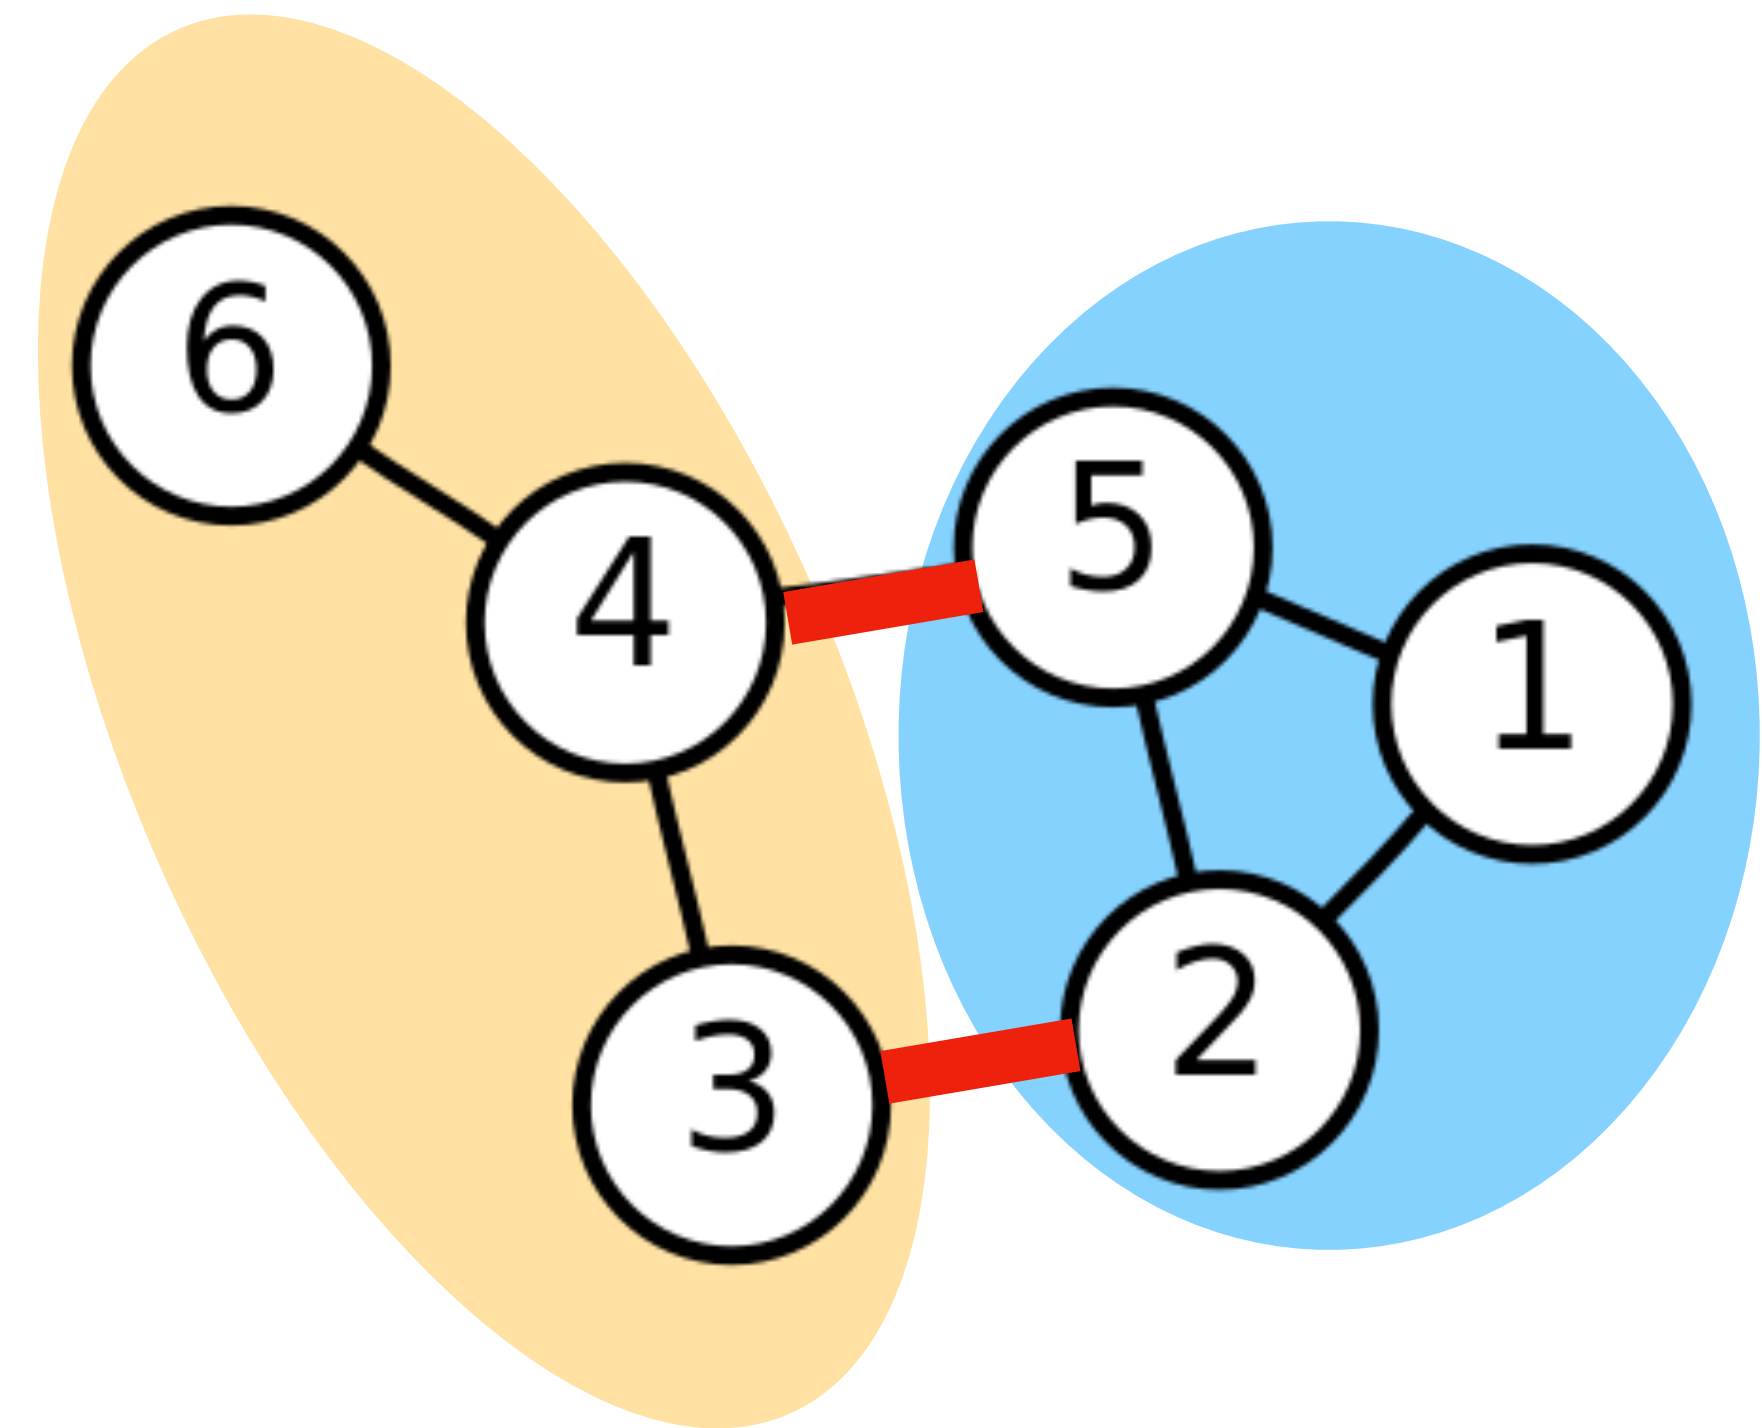
Partition vertices into two disjoint groups A, B

Normalized cut forces clusters to have larger size

$$\text{Ncut}(A,B) = \text{cut}(A,B) \left(\frac{1}{\text{Vol}(A)} + \frac{1}{\text{Vol}(B)} \right)$$

Where $\text{Vol}(A)$ is the sum of the degrees of all of the nodes in A

Then we just $\arg \min_{A,B} \text{Ncut}(A,B)$



Clustering as a bi-partitioning task

Partition vertices into two disjoint groups A, B


How can we define a “good” partition?

OK!

How can we efficiently calculate such a partition?

Not this way... NP-Hard... lets make a “relaxed” solution

Involving eigenvectors of the graph Laplacian matrix



Laplacian matrix [[edit](#)]

Given a [simple graph](#) G with n vertices v_1, \dots, v_n , its Laplacian matrix $L_{n \times n}$ is defined element-wise as^[1]

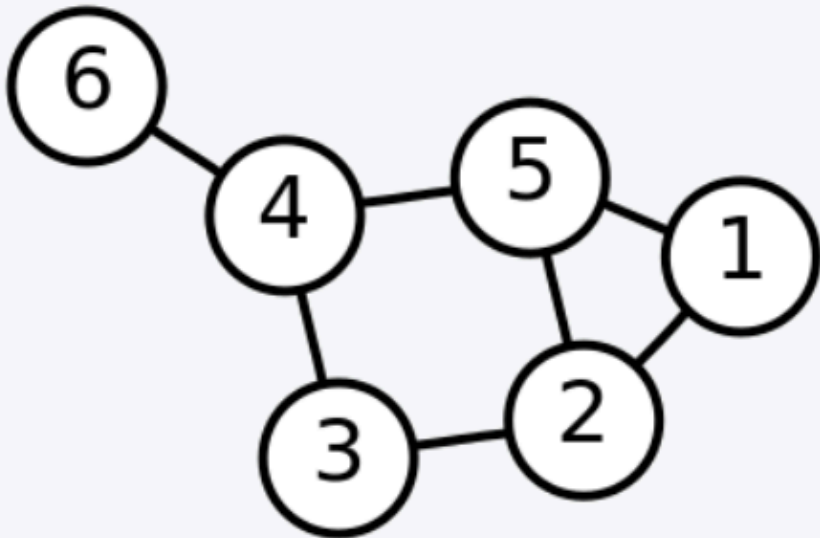
$$L_{i,j} := \begin{cases} \deg(v_i) & \text{if } i = j \\ -1 & \text{if } i \neq j \text{ and } v_i \text{ is adjacent to } v_j \\ 0 & \text{otherwise,} \end{cases}$$

or equivalently by the matrix

$$L = D - A,$$

where D is the [degree matrix](#) and A is the [adjacency matrix](#) of the graph. Since G is a simple graph, A only contains 1s or 0s and its diagonal elements are all 0s.

Here is a simple example of a labelled, undirected graph and its Laplacian matrix.

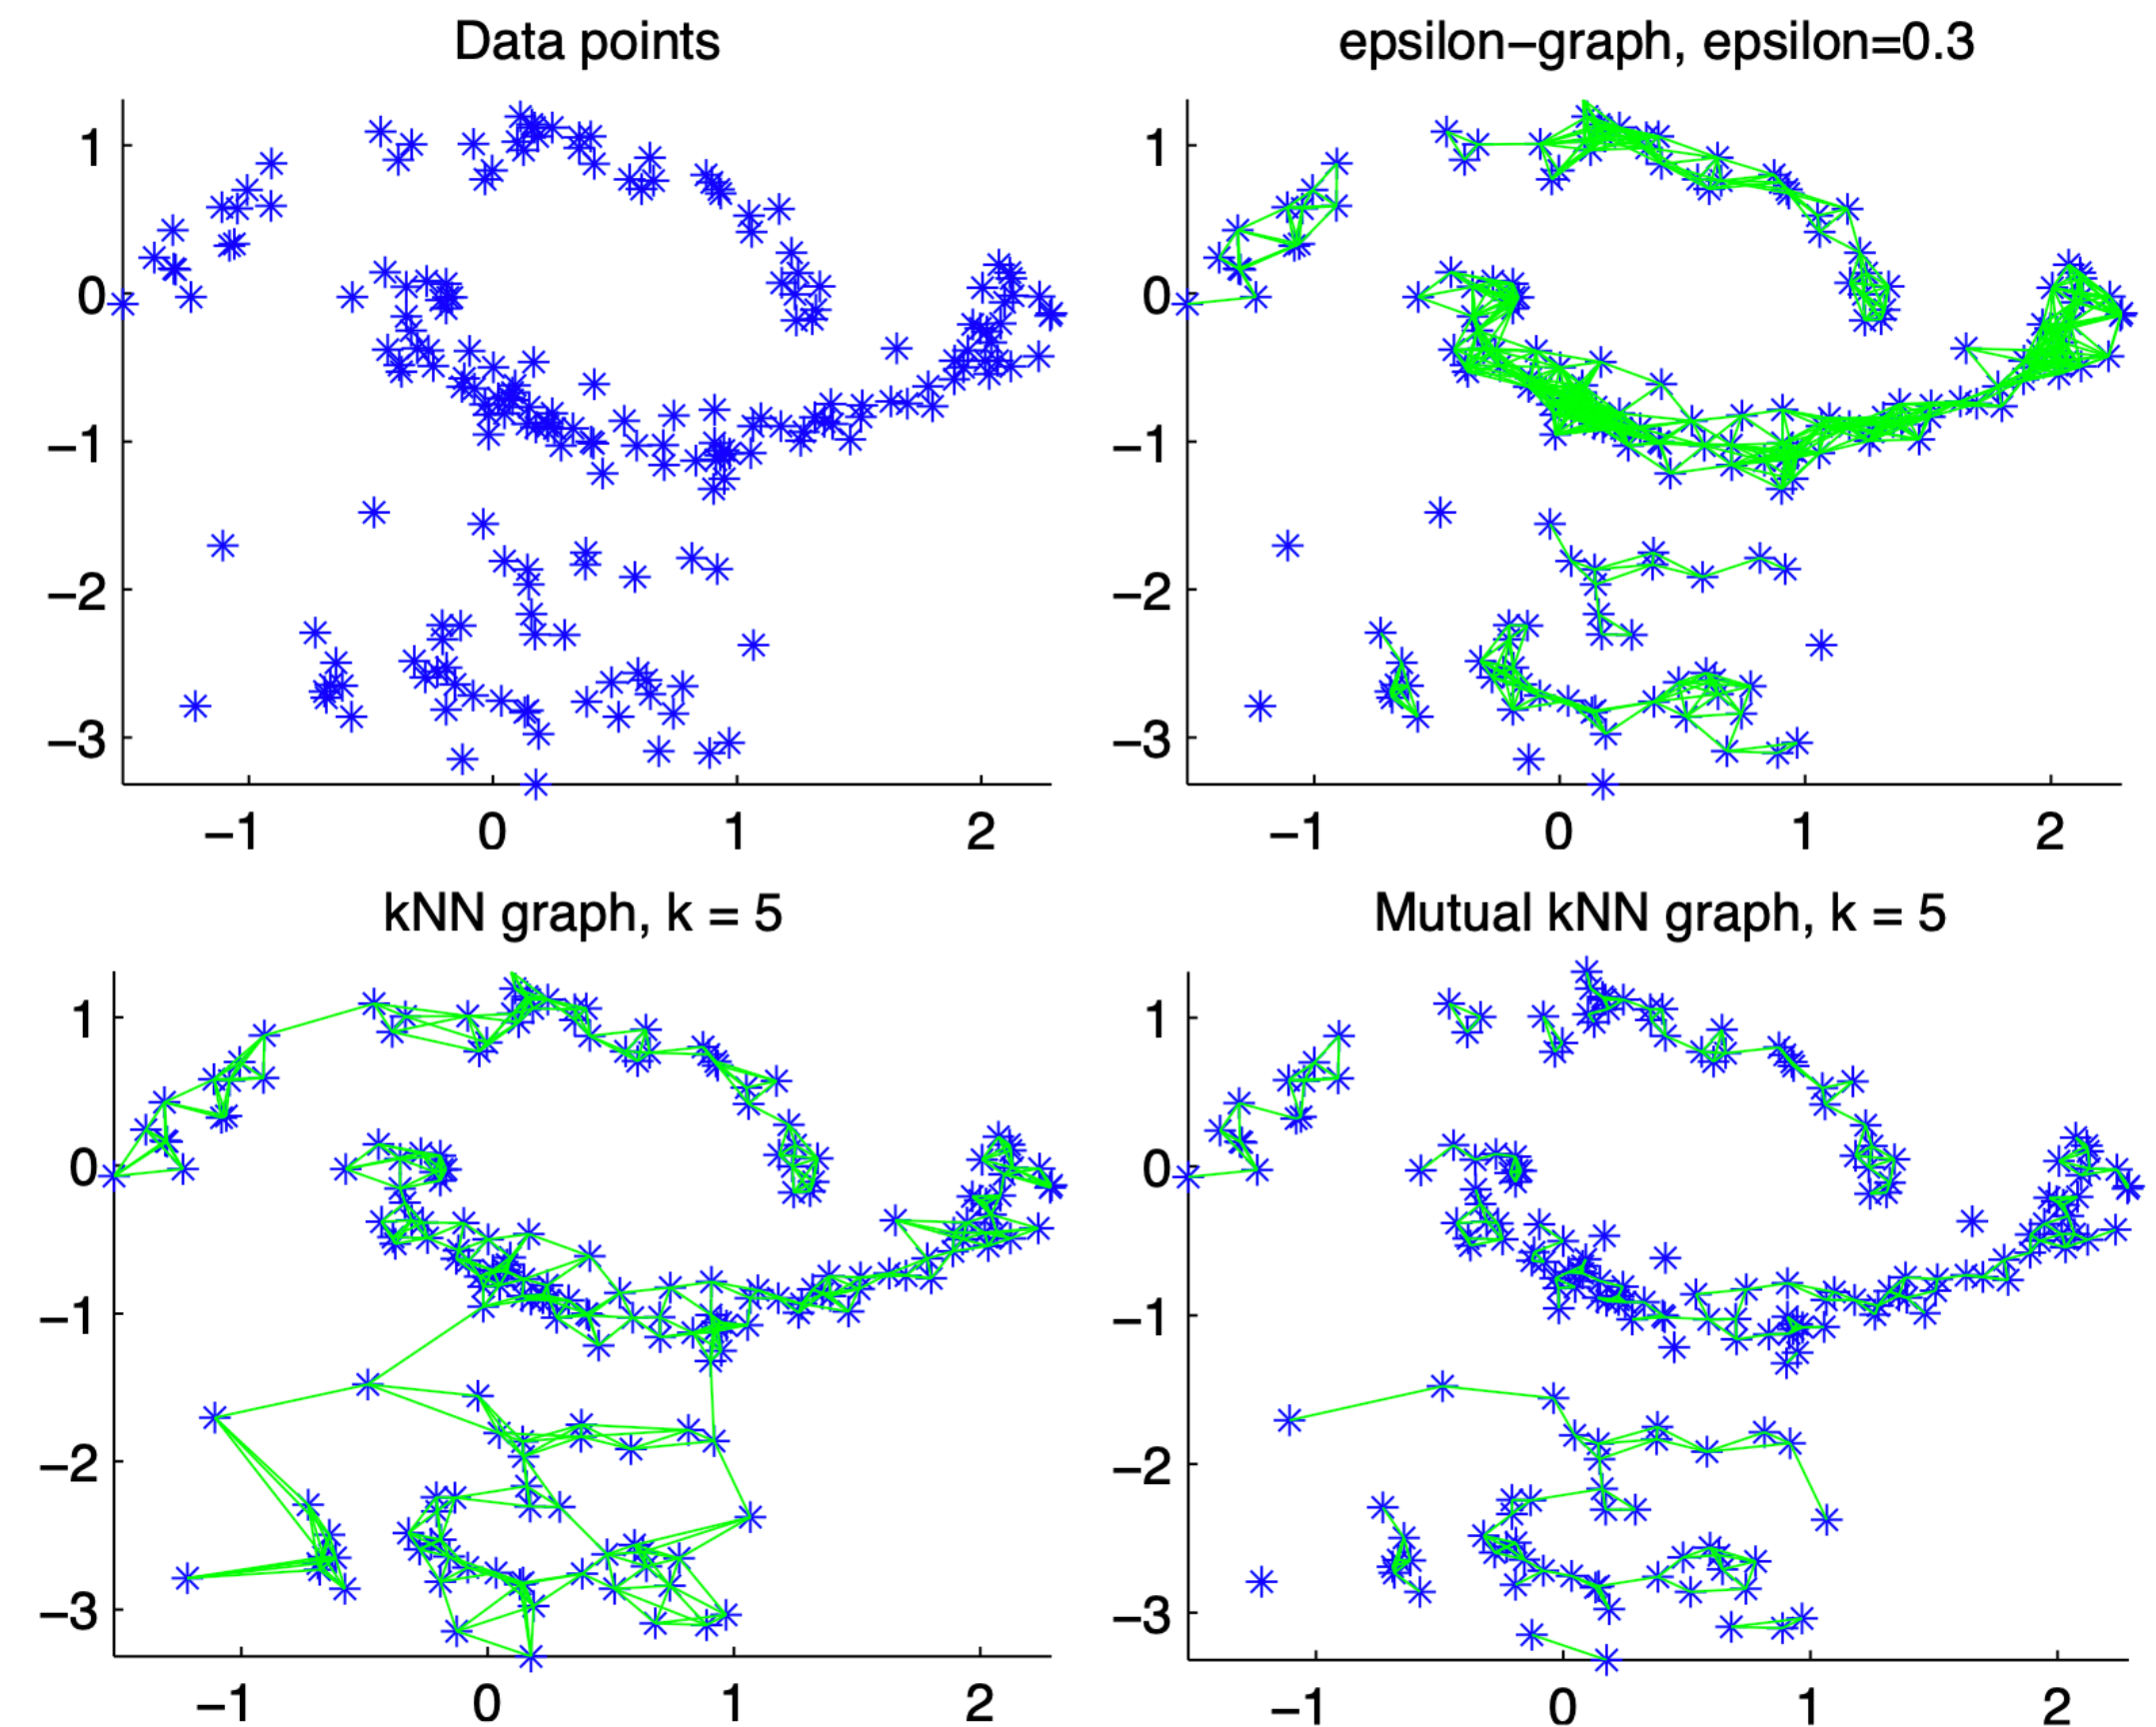
Labelled graph	Degree matrix	Adjacency matrix	Laplacian matrix
	$\begin{pmatrix} 2 & 0 & 0 & 0 & 0 & 0 \\ 0 & 3 & 0 & 0 & 0 & 0 \\ 0 & 0 & 2 & 0 & 0 & 0 \\ 0 & 0 & 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 0 & 3 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 \end{pmatrix}$	$\begin{pmatrix} 0 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \end{pmatrix}$	$\begin{pmatrix} 2 & -1 & 0 & 0 & -1 & 0 \\ -1 & 3 & -1 & 0 & -1 & 0 \\ 0 & -1 & 2 & -1 & 0 & 0 \\ 0 & 0 & -1 & 3 & -1 & -1 \\ -1 & -1 & 0 & -1 & 3 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 \end{pmatrix}$

We observe for the undirected graph that both the [adjacency matrix](#) and the Laplacian matrix are symmetric, and that row- and column-sums of the Laplacian matrix are all zeros.

Lots of different similarity graphs and graph Laplacians!

$$L_{\text{sym}} := D^{-1/2} L D^{-1/2} = I - D^{-1/2} W D^{-1/2}$$

$$L_{\text{rw}} := D^{-1} L = I - D^{-1} W.$$



Normalized Cut and Graph Laplacian

$$\text{Ncut}(A, B) := \text{cut}(A, B) \left(\frac{1}{\text{vol}(A)} + \frac{1}{\text{vol}(B)} \right)$$

$$\text{Let } \mathbf{f} = [f_1 \ f_2 \ \dots \ f_n]^T \text{ with } f_i = \begin{cases} \frac{1}{\text{vol}(A)} & \text{if } i \in A \\ -\frac{1}{\text{vol}(B)} & \text{if } i \in B \end{cases}$$

$$\mathbf{f}^T \mathbf{L} \mathbf{f} = \sum_{i,j} w_{ij} (f_i - f_j)^2 = \sum_{i \in A, j \in B} w_{ij} \left(\frac{1}{\text{vol}(A)} + \frac{1}{\text{vol}(B)} \right)^2$$

$$\mathbf{f}^T \mathbf{D} \mathbf{f} = \sum_j d_i f_i^2 = \sum_{i \in A} \frac{d_i}{\text{vol}(A)^2} + \sum_{j \in B} \frac{d_j}{\text{vol}(B)^2} = \frac{1}{\text{vol}(A)} + \frac{1}{\text{vol}(B)}$$

$$\text{Ncut}(A, B) = \frac{\mathbf{f}^T \mathbf{L} \mathbf{f}}{\mathbf{f}^T \mathbf{D} \mathbf{f}}$$

Normalized Cut and Graph Laplacian

$$\min \text{Ncut}(A, B) = \min \frac{\mathbf{f}^T \mathbf{L} \mathbf{f}}{\mathbf{f}^T \mathbf{D} \mathbf{f}}$$

$$\text{where } \mathbf{f} = [f_1 \ f_2 \ \dots \ f_n]^T \text{ with } f_i = \begin{cases} \frac{1}{\text{vol}(A)} & \text{if } i \in A \\ -\frac{1}{\text{vol}(B)} & \text{if } i \in B \end{cases}$$

$$\text{Relaxation: } \min \frac{\mathbf{f}^T \mathbf{L} \mathbf{f}}{\mathbf{f}^T \mathbf{D} \mathbf{f}} \quad \text{s.t.} \quad \mathbf{f}^T \mathbf{D} \mathbf{1} = 0$$

Solution: \mathbf{f} – second eigenvector of generalized eval problem

$$\mathbf{L} \mathbf{f} = \lambda \mathbf{D} \mathbf{f}$$

Obtain cluster assignments by thresholding \mathbf{f} at 0

Approximation of Normalized cut

$$\text{Ncut}(A, B) := \text{cut}(A, B) \left(\frac{1}{\text{vol}(A)} + \frac{1}{\text{vol}(B)} \right)$$

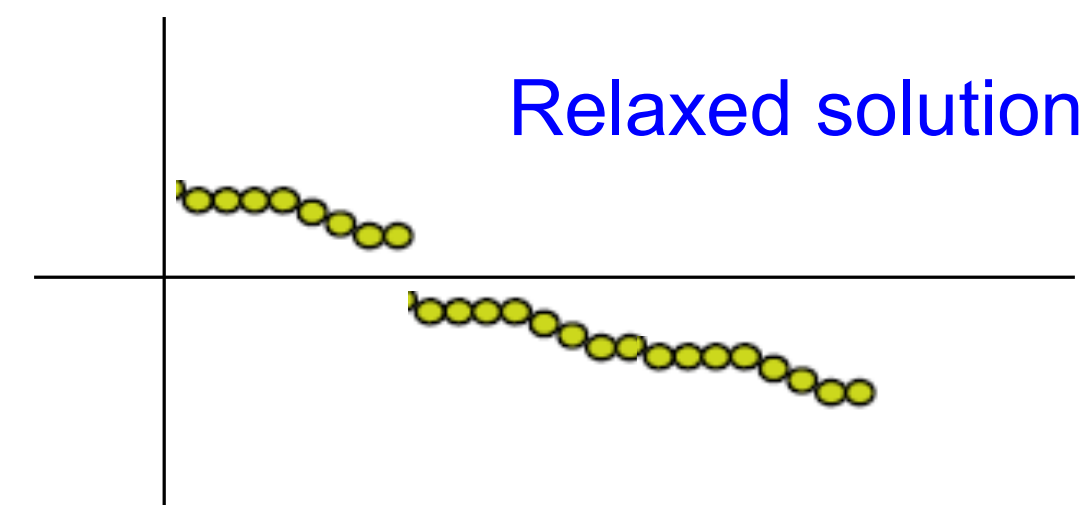

Let f be the eigenvector corresponding to the second smallest eval of the generalized eval problem.

$$\mathbf{L}f = \lambda \mathbf{D}f$$

Equivalent to eigenvector corresponding to the second smallest eval of the normalized Laplacian $L' = D^{-1}L = I - D^{-1}W$



Recover binary partition as follows:

$$\begin{array}{ll} i \in A & \text{if } f_i \geq 0 \\ i \in B & \text{if } f_i < 0 \end{array}$$



How do points cluster?

- Single connected component in graph
 - Binary: sign of 2nd smallest eigenvector (Fiedler eigenvector)
 - Multi: k-Means on 2nd through k+1 eigenvectors
- Multiple connected components in graph (well-separated clusters)
 - Multiplicity of components = multiplicity of 0 eigenvalues and each such eigenvector is a cluster



Spectral clustering

Steps

- Preprocess the data
 - Create a matrix describing connectivity between datapoints (e.g. kNN) or similarity between datapoints (Gaussian kernel)
- Find the graph Laplacian vector of the matrix
 - Several different Laplacians! Each have different properties! Not just one spectral clustering
- Use K-means on evecs/evals of the graph Laplacian to cluster
 - Multiplicity of 0 eigenvalues are the number of connected components in the graph. So if there's no overlap between clusters)... then these smaller evecs assign datapoints to clusters
 - Other options
 - If binary clustering: the 1st (smallest eigenvalue) eigenvector is all 1s, then the sign of the 2nd eigenvector (so-called Fiedler eigenvector) components are the cluster assignments
 - If k clusters: take eigenvectors 2 through k+1 as indicator variables